

NCTU Pattern Recognition, Homework 5

- 309553018 林孟學

Result:

```
Accuracy of my model on test-set: 0.8129
```

Code:

```
175 def test_model(model, test_data):
176     device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
177     y_pred = []
178     for i, (inputs, _) in enumerate(test_data):
179         if i % 5 == 0:
180             print(f'Batch: {i+1}/20')
181             inputs = inputs.to(device)
182             # labels = labels.to(device)
183
184             optimizer.zero_grad()
185
186             outputs = model(inputs)
187             ret, predictions = torch.max(outputs.data, 1)
188             y_pred += [predictions.cpu().numpy()]
189
190     y_pred = np.array(y_pred)
191     y_pred = y_pred.reshape((-1))
192
193
194     ### DO NOT MODIFY CODE BELOW!
195     # please screen shot your results and post it on your report
196
197     assert y_pred.shape == (10000,)
198
199     y_test = np.load("dataset/y_test.npy")
200     print("Accuracy of my model on test-set: ", accuracy_score(y_test, y_pred))
201
202
203     # Train or Test
204     '''
205 > trained_model, history = train_model(...)
211 )
212 '''
213 test_model(resnet50, TestDataLoader)
```

Pretrained Model:

Use pretrained model resnet50.

Modify the last layer of resnet50, and only train these two layers.

```
97 resnet50 = models.resnet50(pretrained=True)
98 for param in resnet50.parameters():
99     param.requires_grad = False
100
101 fc_inputs = resnet50.fc.in_features
102 resnet50.fc = nn.Sequential(
103     nn.Linear(fc_inputs, 256),
104     nn.ReLU(),
105     nn.Dropout(0.4),
106     nn.Linear(256, 10),
107     nn.LogSoftmax(dim=1)
108 )
```

Hyper Parameters:

```
18 BatchSize = 500
19 LearningRate = 0.001
20 EpochCounts = 10
```

Optimizer:

```
111 optimizer = optim.Adam(resnet50.parameters(), lr=LearningRate)
```

Loss function:

```
208 loss_function = nn.NLLLoss(),
loss(input, class) = -input[class]
```

The property we guess of the correct label(input[class]) must be as high as possible →

add minus symbol → as low as possible

Preprocessing:

Just resize and normalize it.

No other transforms such as rotate centerCrop...

```
53 preprocess = transforms.Compose([
54     transforms.ToTensor(),
55     # transforms.Resize(256),
56     # transforms.CenterCrop(224),
57     transforms.Resize(224),
58     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
59 ])
```

Train:

Train model 4 times on the SAME model (Load model after save previous model).

```
215 '''
216     train 1 → Epochs: 5, BatchSize: 500
217     train 2 → Epochs: 5, BatchSize: 500
218     train 3 → Epochs: 10, BatchSize: 500
219     train 4 → Epochs: 10, BatchSize: 500
220 '''
```

GPU:

Use GPU for speed up.

```
116 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```